

## Gradient methods for parameter optimization

### Exercise T4.1: Multilayer perceptron recap

(tutorial)

- (a) Recap the optimization of the MLP parameters (via the backpropagation algorithm).
- (b) Outline the weight space symmetries giving rise to  $\prod_{v=1}^L N_v! \cdot 2^{N_v}$  equivalent solutions where  $L$  is the number of hidden layers and  $N_v$  the respective number of neurons in layer  $v \implies$  no unique global minimum but a large equivalence class of (best) solutions.

### Exercise T4.2: Linear neuron for regression

(tutorial)

To prepare for the homework, we discuss a simple connectionist neuron with linear output function for a real one-dimensional input  $x \in \mathbb{R}$  and output  $y \in \mathbb{R}$ .

- (a) Describe the output function  $y(x; \underline{\mathbf{w}})$  of the neuron in vector notation.
- (b) Derive the gradient and Hessian matrix of the quadratic error function.
- (c) Solve the optimization of the quadratic error function for a data set  $\{(x^{(\alpha)}, y_T^{(\alpha)})\}_{\alpha=1, \dots, p}$  analytically in matrix form.
- (d) Calculate the solution when the objective includes the quadratic training cost  $E^T$  plus a “weight decay” regularization term as used in *ridge regression*, i.e.

$$R_{[\underline{\mathbf{w}}]} = E_{[\underline{\mathbf{w}}]}^T + \lambda \|\underline{\mathbf{w}}\|^2$$

### Exercise T4.3: Conjugate gradient

(tutorial)

- (a) How does the convergence speed of *gradient descent* depend on the learning rate  $\eta$ ?
- (b) Describe how *line search* speeds up convergence.
- (c) What is a *conjugate direction* and how can it improve convergence speed?

### Exercise H4.1: Line search

(homework, 4 points)

In this exercise you will analyze line search based on the simple example of a linear neuron with quadratic cost function  $E_{[\underline{\mathbf{w}}]}^T$ . Here we optimize the cost function along a given direction  $\underline{\mathbf{d}}_t$  (that can be but is not necessarily identical to the gradient  $\underline{\mathbf{g}}_t$ ):

$$\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t - \eta_t \underline{\mathbf{d}}_t.$$

- (a) (1 point) Approximate the cost at the next time step: Derive the 2<sup>nd</sup> order Taylor approximation of an arbitrary cost  $E_{[\underline{\mathbf{w}}_{t+1}]}^T$  around  $\underline{\mathbf{w}}_t$ .

- (b) (1 point) Derive a bound on the step size  $\eta_t$  by using the above approximation in

$$E_{[\mathbf{w}_{t+1}]}^T \stackrel{!}{\leq} E_{[\mathbf{w}_t]}^T.$$

- (c) (1 point) Derive the optimal step size  $\eta_t^*$  for the quadratic cost function

$$E_{[\mathbf{w}]}^T := \frac{1}{2}(\mathbf{w} - \mathbf{w}^*)^\top \mathbf{H}(\mathbf{w} - \mathbf{w}^*)$$

with its minimum at  $\mathbf{w}^*$  by minimizing the cost function w.r.t.  $\eta_t$ . Make sure your solution depends only on known quantities like the weight vector  $\mathbf{w}_t$ , the gradient  $\nabla E_{[\mathbf{w}]}^T|_{\mathbf{w}_t}$  and/or the Hessian  $\mathbf{H}$  of  $E_{[\mathbf{w}_t]}^T$ .

- (d) (1 point) Prove that the gradient  $\nabla E_{[\mathbf{w}]}^T|_{\mathbf{w}_{t+1}}$  after one update step with *line search* is orthogonal to the optimized direction  $\mathbf{d}_t$ .

### Exercise H4.2: Comparison of gradient descent methods (homework, 6 points)

In this exercise we compare the performance of three learning procedures applied to a simple connectionist neuron with a linear output function. All procedures will compute the gradient using the entire training set (batch gradient descent). The procedures are:

- (i) Gradient (or steepest) descent with constant learning rate,
- (ii) steepest descent combined with a line search method to determine the learning rate, and
- (iii) the conjugate gradient method.

**Training Data:** The training data set consists of three samples ( $p = 3$ ):

$$\{(x^{(\alpha)}, y_T^{(\alpha)})\} = \{(-1, -0.1), (0.3, 0.5), (2, 0.5)\},$$

i.e. for a given data point, both input and output are scalar values.

**Cost function:** The gradient for the *quadratic error* function is given by

$$\mathbf{g}(\mathbf{w}) = \frac{\partial E^T}{\partial \mathbf{w}} = \mathbf{H} \mathbf{w} - \mathbf{X} \mathbf{y}_{\text{True}}^\top, \quad \text{with} \quad \mathbf{H} = \mathbf{X} \mathbf{X}^\top,$$

where  $\mathbf{X} = \begin{pmatrix} 1 & 1 & \dots & 1 \\ x^{(1)} & x^{(2)} & \dots & x^{(p)} \end{pmatrix} \in \mathbb{R}^{2,p}$  and  $\mathbf{y}_{\text{True}} = (y_T^{(1)}, y_T^{(2)}, \dots, y_T^{(p)}) \in \mathbb{R}^{1,p}$ .

**Initialization:** Use the following initialization for all three (batch) gradient methods:

$$\mathbf{w}_1 = (w_0, w_1)_1^\top = (-0.45, 0.2)^\top$$

- (a) (2 points) *Gradient Descent:* Implement a steepest descent procedure where the weights at iteration  $t + 1$  are calculated using the weights and the gradient at iteration  $t$

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \mathbf{g}_t,$$

with an adequate learning rate  $\eta$  and where  $\mathbf{g}_t = \mathbf{g}(\mathbf{w}_t)$ . Plot

- (i) the resulting weight vectors from all iterations as a scatter plot ( $w_0$  vs.  $w_1$ ),
- (ii) and ( $w_i$  vs. iterations  $t$ ) in an additional figure,

to show the development of each parameter during gradient descent.

- (b) (2 points) *Line Search*: Implement a line search procedure

$$\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t - \eta \underline{\mathbf{g}}_t, \quad \text{with optimal step size} \quad \eta = \frac{\underline{\mathbf{g}}_t^\top \underline{\mathbf{g}}_t}{\underline{\mathbf{g}}_t^\top \underline{\mathbf{H}} \underline{\mathbf{g}}_t}.$$

Plot the resulting weight vectors from all iterations as

- (i) a scatter plot ( $w_0$  vs.  $w_1$ ),
- (ii) and ( $w_i$  vs. iterations  $t$ ) in an additional figure,

to show the development of the parameters during line search.

- (c) (2 points) *Conjugate Gradient*: Implement a conjugate gradient procedure:

Initialize:  $\underline{\mathbf{w}}_1, \underline{\mathbf{d}}_1 = -\underline{\mathbf{g}}_1$

**while** stopping criterion not satisfied **do**

minimize  $E^T$  along  $\underline{\mathbf{d}}_t$ :  $\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t + \eta_t \underline{\mathbf{d}}_t$  with step size  $\eta_t = -\frac{\underline{\mathbf{d}}_t^\top \underline{\mathbf{g}}_t}{\underline{\mathbf{d}}_t^\top \underline{\mathbf{H}} \underline{\mathbf{d}}_t}$

calculate new gradient  $\underline{\mathbf{g}}_{t+1} = \underline{\mathbf{H}} \underline{\mathbf{w}}_{t+1} - \underline{\mathbf{X}} \underline{\mathbf{y}}_{\text{True}}^\top$

calculate new conjugate direction  $\underline{\mathbf{d}}_{t+1} = \underline{\mathbf{g}}_{t+1} + \beta_t \underline{\mathbf{d}}_t$  with “momentum”

$$\beta_t = -\frac{\underline{\mathbf{g}}_{t+1}^\top \underline{\mathbf{g}}_{t+1}}{\underline{\mathbf{g}}_t^\top \underline{\mathbf{g}}_t}. \quad (\text{Fletcher-Reeves form})$$

increase  $t \leftarrow t + 1$

**end**

Plot the resulting weight vectors from all iterations as

- (i) a scatter plot ( $w_0$  vs.  $w_1$ ),
- (ii) and ( $w_i$  vs. iterations  $t$ ) in an additional figure,

to show the development of the parameters during conjugate gradient descent.

Compare the different methods in terms of convergence behavior.

**Total 10 points.**